

# Sirix - Beyond Versioning of Persistent Tree-Structures

Storing, Querying and Analysing Differences

Johannes Lichtenberger \*

October 10, 2013

## Abstract

Nowadays, disk space is cheap and versioning becomes increasingly interesting. Besides, the days of mechanical disks are numbered and flash drives will certainly predominate in the future. We propose a system which takes full advantage of fast random reads and sequential log-structured writes of flash drives to store fine grained temporal trees. Retaining not only the most recent version but also snapshots of the past facilitates the analysis of temporal patterns.

**Keywords:** Versioning, Git, SVN, Mercurial, Revision, Diffing, XML, JSON, Storage, Visual Analytics

## 1 Introduction

Storage space follows Moore's Law and doubles approximately every two years. It becomes increasingly tempting to store not only the current version, but all past snapshots of data. Hence, it is possible to go back in time in the advent of any kind of failures or unintentional modifications. Furthermore analytical tasks to detect patterns in temporal data are supported, due to the availability of each stored snapshot. Moreover semi-structured data in the form of tree-structures is very common either in the form of XML- or JSON-dialects.

### 1.1 Problem Statement

We face the problem of having to cope with increasingly large data sets due

### 1.2 Contribution

## 2 Conclusion

## 3 Acknowledgements

---

\*email: [lichtenberger.johannes@gmail.com](mailto:lichtenberger.johannes@gmail.com)